

Reliable for Whom?

Directional Reliability in AI-Mediated Political Dialogue

Jaeyoun You · Dept. of Social Innovation Convergence, Hanyang University, Seoul

Reliability in AI mediation is not a property of an output alone, but a directional relationship among speaker, recipient, and system.

The same intervention can be felt as violation or as protection — depending on the direction in which it acts.

1 The problem: symmetric harmlessness

ML evaluation scores harmlessness as a symmetric property of text — a toxicity value on an output. When AI mediates human-to-human dialogue, mechanically removing all friction manufactures a sanitized consensus and erases the conflict pluralism requires. In polarized settings, fact-checks alone often deepen defensiveness rather than restore engagement.

2 Conceptual move (Mouffe)

Agonistic pluralism transforms antagonism (enemies) into agonism (adversaries who respect each other's right to speak) — rather than eliminating conflict. The deeper crisis is exclusion: people made to feel their voice no longer counts.

A mediating AI can act as a provisional communicative scaffold: temporary support that lowers affective barriers to continued interaction without resolving the underlying dispute.


3 Directional Reliability

Users treat outgoing speech as expressive sovereignty, yet welcome buffering of what they receive. Legitimacy tracks the direction of the intervention, not the content of the text.

DIRECTIONALLY RELIABLE =

preserve the sender's control over expression while offering the recipient optional receptive scaffolding. Imposed buffering reproduces paternalism — optionality is essential.

WHERE THE INTERVENTION ACTS



Entrenched users reject both; reluctant users accept recipient-side scaffolding as re-entry support. A reliable scaffold is also temporary — it should fade as unmediated exchange stabilizes.

4 Design probe

N = 25 · exploratory, qualitative · IRB-approved

Real-time chat in which an LLM rewrote each message by conflict style — accommodating, collaborative, competing, compromising — while preserving intent.

Run May 2025, after South Korea's Dec 3 2024 martial-law declaration. Focal issue: election-fraud allegations — epistemic deadlock where factual correction no longer restores engagement.

5 What the probe surfaced

TWO PARTICIPANT GROUPS

Politically entrenched

Resist intervention in BOTH directions. The raw, unmediated exchange is itself the authenticity they seek.

Reluctant participants

Hold views but had stopped speaking. Reception-side scaffolding made participation feel manageable — a re-entry support, not a smoothing-over.

Among reluctant users, self-reported willingness to continue was higher under accommodating than competing mediation (descriptive: $M \approx 5.8$ vs 3.2 on 7 pts, $n \approx 6$ /cell; no significance test — interviews carry the weight).

↓ THE PATTERN (synthesis)

Expression = sovereignty (editing felt as violation). Reception = welcome buffering (felt as protection). Legitimacy tracks direction, not content.

6 Implications for ML evaluation

1. Asymmetric objectives

Penalize distortion of sender intent; credit recipient-side buffering that curbs escalation without overriding agency.

2. Relational capability metrics

Expressive fidelity, perceived agency, willingness to continue, reciprocal turn-taking — disaggregated by user type.

3. Dynamic fading

A scaffold should be temporary: fade buffering as users' unmediated drafts grow less hostile.

SCOPE & LIMITS

An exploratory, hypothesis-generating probe. Whether directional reliability generalizes across other polarization structures remains open.

TAKEAWAY

Ask not only whether an output is safer — but for whom, in which direction, and at what cost to agency.

