
Reliable for Whom? Directional Reliability in AI-Mediated Political Dialogue

Jaeyoun You¹

Abstract

Many ML evaluation frameworks approach reliable AI through epistemic accuracy, robustness, and general harmlessness. Yet as generative AI systems increasingly mediate human-to-human communication, these model-centric notions become insufficient. This paper is an empirically grounded position paper. Drawing on a formative design probe ($N=25$) conducted in the aftermath of South Korea’s martial law declaration (December 3, 2024), we argue that reliability in AI mediation is better understood as a *directional relationship* among speaker, recipient, and system than as a property of a model output considered in isolation. The probe suggested a layered asymmetry: politically entrenched participants resisted AI intervention in both directions, while participants who held views but had been reluctant to voice them reported that reception-side scaffolding appeared to lower the threshold for participation. We offer *Directional Reliability* as an orienting analytic concept developed from the probe: rather than treating harmlessness as a symmetric property of text, it asks whether an intervention preserves expressive agency on the sender side while offering optional scaffolding for reception on the recipient side. This reframing surfaces a blind spot in current output-level evaluation frameworks for AI systems deployed in politically difficult dialogue. We report the empirical material as hypothesis-generating rather than confirmatory; the contribution is conceptual.

1. Introduction: The Problem with Symmetric Harmlessness

Many ML evaluation frameworks for trustworthy AI emphasize epistemic fidelity, robustness, and safety. Similarly, contemporary alignment discussions often focus on objectives such as helpfulness, honesty, harmlessness, and related safety goals (Askell et al., 2021). Within these frameworks, reliability is commonly evaluated through factual accuracy, robustness, and the reduction of harmful or toxic outputs. While appropriate for systems whose primary role is in-

formation provision, these criteria may be insufficient for AI systems that actively mediate human-to-human communication. In polarized political environments, fact-checks alone may fail to restore engagement and can provoke defensive reactions. When safety filters mechanically remove all conflict to ensure “harmlessness,” they risk producing a sanitized consensus that obscures the friction inherent in pluralistic dialogue.

To fix terms, we use *harmlessness* to mean the absence of toxic or offensive content in an output, *safety* to mean the avoidance of downstream harm from system behavior, and *reliability* to mean the degree to which users can depend on a system to perform its intended role; *trustworthiness* we treat as the broader disposition these properties support. Our argument concerns reliability specifically: for mediating systems, reliability has a directional structure that output-level harmlessness does not capture. We do not propose to redefine the other terms.

This gap has become increasingly consequential as generative AI moves from passive information retrieval toward active participation in human-to-human communication. AI systems now summarize, rephrase, moderate, and reframe messages in real time—intervening not just in what is said, but in how it is received and by whom. In these settings, the relevant question is no longer only whether a model produces accurate or inoffensive outputs, but whether its interventions are experienced as legitimate by the people whose speech it shapes.

Reliability in human-to-human mediation is not about producing polite text. It is about *who* is intervening, in *which direction*, and with *what authority*. This paper does not propose a general theory of AI reliability. Through a small design probe, it surfaces a blind spot: the same AI intervention can be experienced as violation or as protection depending on the communicative direction in which it acts. The central claim is that reliability in AI-mediated political dialogue can be *experienced directionally*. Existing evaluation frameworks typically assess interventions at the level of outputs, yet our observations suggest that users may evaluate the same intervention differently depending on whether it acts upon expression or reception. This asymmetry deserves a place in ML evaluation. We introduce **Directional Reliability** as an orienting concept for making it visible.

2. Conceptual Move: Why Harmlessness Is Not Enough

Many ML safety frameworks assume that filtering emotional toxicity will allow rational consensus to emerge, broadly consistent with Habermasian ideals of communicative rationality (Habermas, 1984). Chantal Mouffe (Mouffe, 1999) offers a normative warning against this. In her framework of agonistic pluralism, democratic engagement aims to transform *antagonism* (a struggle between enemies) into *agonism* (a struggle between adversaries who respect each other’s right to speak), rather than to eliminate conflict. An AI that systematically sanitizes political expression does not resolve conflict; it may suppress the friction that democratic pluralism requires.

Mouffe’s later work sharpens this further. In *For a Left Populism* (Mouffe, 2018), she argues that the crisis of liberal democracy stems less from an excess of conflict than from the exclusion of large portions of the population from meaningful participation. On her account, people have been made to feel that their voices do not count. This resonates with our empirical findings: the participants who benefited most from AI scaffolding were not the politically entrenched, but those who had opinions yet had stopped feeling able to express them.

We draw on Mouffe as a normative check, not a design blueprint. Her account tells us what an AI mediator should *not* do (erase conflict wholesale), and points toward what democratic participation actually requires (conditions for those effectively silenced to speak). We propose that in conditions of deep polarization, a mediating AI can serve as a *provisional communicative scaffold*: a temporary support that lowers affective barriers to continued interaction without resolving the underlying dispute.

Mouffe does not distinguish between sender-side and recipient-side interventions. However, her account highlights a broader concern: democratic participation depends not only on the content of communication but also on the conditions under which people are able to enter and sustain political dialogue. This concern motivates our later focus on the directionality of AI-mediated interventions.

3. Design Probe: Observing Directional Asymmetry

To explore how such scaffolding operates in practice, we conducted a concept-generating design probe ($N=25$ human participants). We built a real-time chat interface where an LLM rewrote users’ messages according to conflict-resolution styles (accommodating, collaborative, competing, compromising) while preserving semantic intent. The probe was not designed to validate a general theory but to gener-

ate conceptual insights about how users negotiate trust and agency when AI intervenes in their communication.

Methodological Context. The study was conducted in May 2025, in the aftermath of President Yoon Suk-yeol’s declaration of martial law on December 3, 2024, and the civic resistance and impeachment proceedings that followed. The focal issue, electoral fraud allegations, illustrates epistemic deadlock at its most acute. A significant segment of the conservative electorate had come to believe that the 2020 legislative elections were fraudulent; others rejected this entirely. These opposing convictions, each held with high affective intensity, contributed to deepening partisan polarization—a dynamic President Yoon himself cited as justification for the martial law declaration. In this context, factual correction alone was insufficient to restore mutual engagement; the dispute had long ceased to be primarily about evidence.

Participants (11 aged 18–29; 14 aged 30–49) were recruited via university communities and social media. Of the 25, 4 (16%) believed in electoral fraud and 21 (84%) did not, roughly reflecting national survey distributions at the time. Given the uneven natural distribution, establishing human-to-human dyads for all sessions was not feasible. To balance conditions, one AI persona representing a pro-fraud stance was introduced as a conversation partner; crucially, participants engaged under the belief that they were communicating with another human peer, preserving the ecological validity of human-to-human mediation. All quantitative analyses focus exclusively on the 25 human participants. Each session included a 20-minute structured chat, post-survey, and semi-structured interview. The study was IRB-approved; all participants gave informed consent and were fully debriefed.

We treat this probe as exploratory and primarily qualitative. The four conditions were not randomized for causal inference, per-condition cells are very small (approximately six participants each), and the study was not pre-registered; we therefore conducted no significance testing, and none should be inferred from the descriptive figures we report. With a sample of this size and multiple survey items, any isolated group difference would be uninterpretable as a statistical effect, so we report survey responses only to convey the direction of a pattern whose primary evidence is qualitative. We retain the raw figures rather than omit them, in the interest of full reporting. This sample is also analytically relevant to the paper’s concern with uneven participation: the probe does not model an ideal deliberative public, but a politically asymmetric setting in which some participants are more willing and able to speak than others. Such affective polarization is globally widespread, from election denialism in the US and Brazil to worldwide vaccine discourse (McCoy et al., 2018), making this a theoretically useful stress test.

The probe surfaced three layered findings:

3.1. The Politically Entrenched: Resistance in Both Directions

Participants with strong prior political commitments resisted AI intervention in both directions. Expression-side editing was rejected as inauthentic (“my voice was lost”), but reception-side filtering was also sometimes viewed with suspicion, as if concealing the opponent’s “true face.” For these participants, the raw unmediated exchange was itself the sought authenticity. AI scaffolding, regardless of direction, felt like interference with the political stakes they were invested in.

3.2. The Reluctant Participants: Scaffolding as Re-Entry Support

A markedly different pattern emerged among participants who held political views but had avoided debate out of fear of attack, embarrassment, or a sense of futility. For this group, which can be interpreted through Mouffe’s account of democratic exclusion, reception-side scaffolding did not merely make the conversation feel more comfortable; participants reported that it made participation feel more manageable. As an illustration of this direction, self-reported willingness to engage in future political discussion was higher, on average, in the accommodating condition ($M=5.8$ on a 7-point scale, $n \approx 6$) than in the competing condition ($M=3.2$, $n \approx 6$). We report these means for transparency rather than as estimates of an effect: given the small cells and the absence of randomization or significance testing, they offer descriptive context for the interview accounts rather than evidence of a generalizable difference. The AI scaffold did not make adversaries like each other; rather, it appeared to temporarily restore some of the interactional conditions under which reluctant participants felt able to speak.

3.3. The Underlying Pattern: Directional Asymmetry

Across both groups, a consistent directional asymmetry emerged. Users treated their outgoing utterances as a domain of expressive sovereignty; expression-side editing was experienced as a violation of agency and authenticity. Yet the same users, especially reluctant participants, welcomed recipient-side buffering as protective rather than intrusive. The legitimacy of AI intervention depended not on the content of the text but on the communicative direction in which the AI acted.

4. The Core Insight: Directional Reliability

We use the term *Directional Reliability* to describe whether users experience an AI mediator as dependable within a

particular communicative role, rather than to denote the objective correctness of an intervention. Existing evaluation frameworks often assess interventions at the level of outputs—for example, whether a response is accurate, safe, or non-toxic. Our probe suggests that users may also evaluate interventions relationally. The same intervention can be experienced differently depending on whether it acts upon a speaker’s expression or a recipient’s reception. From this perspective, reliability in AI-mediated communication is not only a property of outputs but also a property of how interventions are situated within a communicative relationship. This does not imply that conventional notions of reliability are unimportant; rather, it suggests that output-level measures alone may be insufficient for evaluating AI systems that actively mediate human-to-human interaction.

Within this framework, we offer **Directional Reliability** as an orienting analytic concept, developed abductively from the asymmetry observed in our probe rather than deduced from political theory. We avoid the label “sensitizing concept” in its strict methodological sense, since the construct is our analytic term rather than an in-vivo expression drawn from participants’ own vocabulary; its role is to orient evaluation design toward directionality. A mediating AI may be understood as directionally reliable when it preserves the sender’s control over their own expression while offering recipients optional forms of receptive scaffolding. Optionality matters here: if reception-side scaffolding is imposed rather than chosen, it risks reproducing the same paternalism that expression-side editing can create for speakers. Mouffe’s framework helps explain why the wholesale removal of expressive friction may be normatively problematic. The specific asymmetry we identify, however, emerges from how participants negotiated trust with AI mediators in practice. For reluctant participants, reception-side scaffolding did not remove political disagreement; rather, it appeared to help restore some of the conditions under which they felt able to enter and remain in difficult political dialogue.

5. Implications for ML Evaluation

If reliability can be experienced directionally in AI-mediated communication, evaluation frameworks for socially situated AI systems may need to account for this asymmetry. A static toxicity score is a limited basis for evaluating a mediating AI, because the legitimacy of an intervention depends on where it occurs in the communicative relation, not only on what the output contains.

1. **Asymmetric Evaluation Objectives.** Rather than treating toxicity as a universal property of a string, mediation systems should distinguish between two design objectives: preserving expressive fidelity for the sender and providing receptive buffering for the recip-

ient. Evaluation could penalize distortions of sender intent while recognizing recipient-side buffers that reduce destructive escalation without overriding speaker agency.

2. **Relational Capability Metrics.** Following a capability-oriented perspective (Sen, 2009), the relevant question is not only whether the model produces safer text, but whether the system expands users' practical capability to participate in difficult dialogue. Evaluation criteria should include expressive fidelity, perceived agency, willingness to continue interaction, and reciprocal turn-taking, disaggregated by user type. Gains in participation willingness among reluctant users may be more significant than aggregate averages suggest.
3. **Dynamic Fading.** A provisional scaffold should be temporary. A reliable mediating AI may be evaluated on its capacity for *dynamic fading*: gradually reducing receptive buffering as interlocutors build sufficient mutual engagement to sustain less-mediated exchange. One proxy for this threshold could be a sustained reduction in hostility within users' raw, unmediated drafts—a signal that the affective conditions for unscaffolded agonism are beginning to stabilize. Whether directional reliability generalizes across different polarization structures remains an open empirical question.

A further implication concerns alignment itself. In politically charged dialogue, notions of harmlessness and safety are often contested: what counts as “harmless” may depend on whose expression is modified, whose affective burden is reduced, and whose terms of engagement are privileged. Directional Reliability helps make this asymmetry visible at the interactional level, asking not only whether an output is safer, but for whom, in which direction, and at what cost to agency.

6. Conclusion

What counts as trustworthy? When generative AI mediates human conflict, trustworthiness may require more than producing harmless text. It may require recognizing the directional nature of human trust. Reconceptualizing reliability as a directional capacity shifts the question from whether a model produces harmless text to whether an intervention preserves the relational conditions under which difficult speech can continue, and for whom.

We do not claim that reliability is always directional. Our probe suggests that reliability may be experienced directionally in AI-mediated political dialogue, and that this asymmetry deserves consideration in ML evaluation. The finding that most concerns us is the specific pattern among

reluctant participants: people who held political views but had stopped feeling able to express them. For this group, reception-side scaffolding did not appear to smooth over democratic conflict; rather, it appeared to reopen a door to participation that many participants felt had effectively closed. Whether this constitutes a meaningful restoration of civic voice, or merely a more manageable way of engaging with disagreement, remains a normative question that the ML community should not leave unanswered.

References

- Amanda Askell, Yuntao Bai, Anna Chen, Dawn Drain, Deep Ganguli, Tom Henighan, Andy Jones, Nicholas Joseph, Benjamin Mann, Nova DasSarma, Nelson Elhage, Zac Hatfield-Dodds, Danny Hernandez, Jackson Kernion, Kamal Ndousse, Catherine Olsson, Dario Amodei, Tom B. Brown, Jack Clark, Sam McCandlish, Chris Olah, and Jared Kaplan. A general language assistant as a laboratory for alignment. *CoRR*, abs/2112.00861, 2021. URL <https://arxiv.org/abs/2112.00861>.
- Jürgen Habermas. *The Theory of Communicative Action, Vol. 1: Reason and the Rationalization of Society*. Beacon Press, 1984.
- Jennifer McCoy, Tahmina Rahman, and Murat Somer. Polarization and the global crisis of democracy: Common patterns, dynamics, and pernicious consequences for democratic polities. *American Behavioral Scientist*, 62(1): 16–42, 2018.
- Chantal Mouffe. Deliberative democracy or agonistic pluralism? *Social Research*, 66(3):745–758, 1999.
- Chantal Mouffe. *For a Left Populism*. Verso, 2018.
- Amartya Sen. *The Idea of Justice*. Harvard University Press, 2009.